Sahil Khairnar

(+91) 9370070202 | sahilk2303@gmail.com | linkedin/sahil-khairnar | github/sahilk23

Education

University of Mumbai
S.I.E.S Graduate School of Technology — *Navi Mumbai*
B.E in Artificial Intelligence and Data Science — *2022 – 2026*

Technical Skills

AI & LLM: AI Agent Development, LLMs, Prompt Engineering, RAG, LangChain, LangGraph, PEFT, QLoRA, LoRA, Hugging Face Transformers.
Enterprise AI Platforms: Microsoft Copilot Studio, Zapier, OpenAI API, Gemini API.
Documentation & KB Management: Markdown, Knowledge Base (KB) Management, Version Logs, Change Records, User Guides, Test Case Writing.
Backend & APIs: Python, SQL, FastAPI, JSON, Docker Compose.
Databases: PostgreSQL, MongoDB, Pinecone (Vector DB).
Data & ML: PyTorch, TensorFlow, Scikit-Learn, Model evaluation, NumPy, Pandas, Matplotlib, Seaborn, Plotly, Power BI, Tableau.
Developer Tools: Git, GitHub, Docker, Jira, Microsoft 365 (Teams, SharePoint, OneDrive), Jupyter Notebook.

Work Experience

Artificial Intelligence Intern — Dec '25 – Mar '26
Vizpay Business Solution Pvt. Ltd. — *Thane (On-site)*

- Developed and deployed an AI customer support assistant for the internal team enabling natural language querying of a PostgreSQL database; designed a dual-adapter pipeline (intent classifier + NL2SQL) to handle structured enterprise queries end-to-end.
- Fine-tuned Qwen2.5-Coder 14B via QLoRA (PEFT) on 15,000+ multi-table SQL training records; Executed structured test cases and evaluated response quality on a held-out test set, achieving 94% result-set accuracy, documented results, flagged failure modes, and iterated on prompt engineering to close gaps.
- Architected and containerized the full Text-to-SQL pipeline for on-premise deployment, reducing manual SQL execution dependency of the support team by 30%; maintained deployment documentation, version logs, and configuration change records throughout.
- Built a Retrieval-Augmented Generation (RAG) knowledge base pipeline using Pinecone and LangChain over 20+ internal company documents; organized and formatted knowledge base (KB) files to support accurate, context-specific responses to generalized customer support queries.

Data Engineering Intern — Jul '25 – Aug '25
White Warrior — *Thane (On-site)*

- Automated data extraction and cleaning from 50+ web sources using RPA workflows; structured outputs for business analysis and maintained organized, formatted data documentation for team review.
- Triaged and logged operational feedback from field teams using Jira; prepared structured update summaries for stakeholder review and ensured timely resolution across workflow deliverables.

Projects

AI Agent – Hospital Management Assistant

- Developed a multi-agent healthcare automation platform (Python, Gemini API, SQLite) that streamlined the full patient journey from intake to triage, report parsing, appointment and booking, reducing manual coordination effort by 40-60% across front-desk workflows. Authored role-specific user guides and CLI workflow documentation.
- Built role-based Staff/Doctor CLI workflows with natural-language clinical queries, improving doctor access to schedules and patient history and cutting information retrieval time from 10-15 mins to <1 mins.
- Implemented asynchronous orchestration for parallel triage + report processing, decreasing case processing turnaround by 30-45% and enabling faster routing to Cardiology/Neurology/Orthopedics/General departments.

Customer Churn Prediction

- Among the four models evaluated, Random Forest achieved the highest accuracy at 96.60%, with Decision Trees and SVM also performing well at 94.44% and 90.72%, respectively.
- The AUC of the Random Forest model is 0.987, indicating the model is able to distinguish between the two classes (e.g., retained and churned) exceptionally well.